

Voltage Park — Single Node Text-to-Video API

Overview

This project implements a **Kubernetes-deployed asynchronous text-to-video API** using **FastAPI**, **Celery**, **Redis**, and GPU-bound **Rust workers** orchestrated by Python.

It satisfies the Voltage Park take-home assignment for a single node with **8× H100 GPUs**.

Core idea:

- **FastAPI** serves as the HTTP API gateway.
- **Celery** manages job orchestration and retries.
- **Redis** acts as the broker/result backend and status cache.
- **Rust GPU workers** run the `genmo/mochi-1-preview` model.
- **PVC** provides fast local scratch space for video generation.
- **MinIO** stores completed videos for retrieval via presigned URLs.

High-Level Architecture

-  arch_diagram

Components

1. FastAPI — API Gateway

- Async, high-performance Python web framework.
- Handles:
 - **POST /jobs** → Create a job, return `job_id`
 - **GET /jobs/{id}** → Status (`pending`, `processing`, `completed`, `failed`)
 - **GET /jobs** → Paginated job list
 - **GET /jobs/{id}/result** → Presigned download URL (MinIO) or direct MP4

Scaling:

- Stateless → scale horizontally (HPA) by CPU/memory or RPS.
- No GPU allocation.

2. Celery — Job Orchestration

- Mature Python task queue framework.
- Handles retries, routing, and concurrency limits.
- Queue separation:
 - `short-low` — quick jobs
 - `long-high` — heavy/long-running jobs

Scaling:

- HPA based on `queue_depth`.
 - Workers pinned to GPU resources via Kubernetes.
-

3. Redis — Broker + Cache

- Fast, in-memory datastore.
- Stores:
 - Celery task queue
 - Task results
 - Job status cache

Scaling:

- Single instance with AOF persistence for MVP.
 - Optional Redis Cluster for high concurrency.
-

4. Rust GPU Workers

- Long-running processes in GPU pods.
- Host the mochi-1 model for text-to-video.
- Python task (Celery) calls Rust via **gRPC** (preferred) or **PyO3 FFI**.

GPU Binding:

```
resources:
  limits:
    nvidia.com/gpu: 2
```

- Each worker pod processes tasks sequentially per GPU to avoid memory contention.

5. PVC — Persistent Scratch Storage

- **Local NVMe-backed** (or fast network) persistent volume for:
 - Temporary frames
 - Partial encodes
 - Prevents data loss on pod restarts.
-

6. MinIO — Artifact Storage

- **S3-compatible** object store deployed in the cluster.
 - Stores final videos.
 - API returns **presigned URLs** so frontend can download directly without routing through the API pods.
-

gRPC API Specification (Python ↔ Rust Workers)

The Rust workers expose a **gRPC service** consumed by the Python Celery tasks. This keeps orchestration in Python while GPU-heavy work stays in Rust.

Proto file: **video_generator.proto**

```
syntax = "proto3";

package video;

service VideoGenerator {
    // Submit a text-to-video generation request
    rpc Generate (GenerateRequest) returns (GenerateResponse);

    // Check job status (optional if Celery handles status)
    rpc GetStatus (StatusRequest) returns (StatusResponse);
}

message GenerateRequest {
    string job_id = 1;           // UUID from Celery
    string prompt = 2;          // User's text prompt
    int32 resolution_width = 3; // e.g., 1920
    int32 resolution_height = 4; // e.g., 1080
    int32 duration_seconds = 5;  // Clip length
    string quality = 6;          // "low", "medium", "high"
}

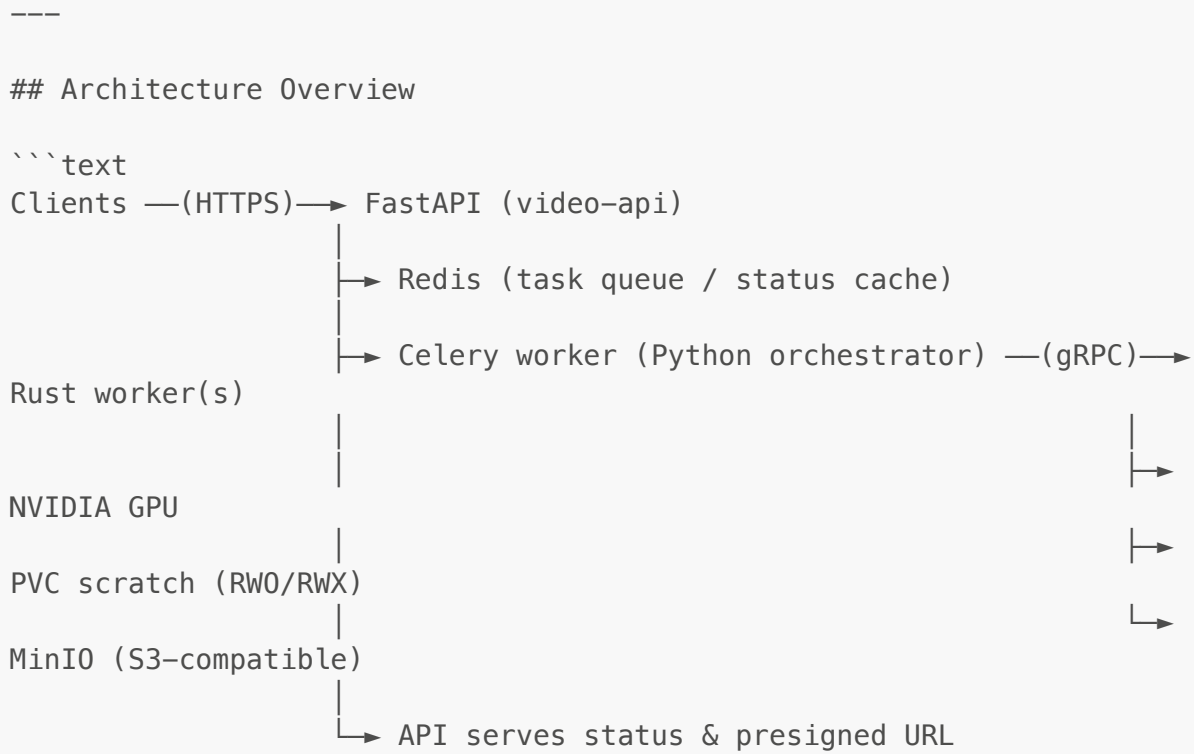
message GenerateResponse {
    bool accepted = 1;          // Whether the worker accepted the job
    string message = 2;         // Info or reason for rejection
}

message StatusRequest {
    string job_id = 1;
}

message StatusResponse {
    string job_id = 1;
    string state = 2;           // "pending", "processing",
                                // "completed", "failed"
    string output_path = 3;     // Path on PVC or S3 URL
    string error_message = 4;   // If failed
}

# gRPC API & Kubernetes Deployment Specification (Python ↔ Rust Workers)

**Goal:** Keep orchestration in Python (FastAPI + Celery) while
offloading GPU-heavy work to Rust workers exposed via gRPC. Workers
run in the same pod as the GPU process (or alone) and write
intermediates to a PVC, then upload final artifacts to MinIO.
```



Flow:

1. FastAPI enqueues a Celery job with **job_id** and request payload.
2. Celery task dials the Rust worker **gRPC** service (**Generate**) inside the cluster (same namespace).
3. Rust worker loads the model (on first call), generates video to a **PVC scratch** path.
4. Worker uploads final video to **MinIO** and records job state (Redis update or callback to API).
5. API exposes status + **presigned URL** to clients.

Advantages:

- **Language isolation:** Rust can evolve independently of orchestration.
- **Clear contracts:** gRPC schema is the source of truth.
- **Scales cleanly:** Add workers; API/Celery remain unchanged.

gRPC API

.proto (tonic-compatible)

```

syntax = "proto3";
package video;

service Generator {
  rpc Generate(GenerateRequest) returns (stream GenerateEvent);
  rpc GetJob(GetJobRequest) returns (GetJobResponse);
}

```

```

message GenerateRequest {
    string job_id = 1;           // Idempotency key
    string prompt = 2;          // Text or JSON-encoded params
    string model = 3;           // e.g. "mochi-1-preview"
    string scratch_dir = 4;      // Mounted PVC path (pod-local)
    map<string,string> options = 5; // width, height, fps, seed,
    duration, etc.
}

message GenerateEvent {
    oneof event {
        Progress progress = 1;
        Result result = 2;
        Error error = 3;
    }
}

message Progress {
    string job_id = 1;
    int32 step = 2;
    int32 total_steps = 3;
    string message = 4;
    double gpu_util = 5;        // Optional, for live UX/metrics
}

message Result {
    string job_id = 1;
    string artifact_path = 2;    // Final local path before upload
    string s3_url = 3;           // s3://bucket/key
    string etag = 4;
    int64 bytes = 5;
}

message Error {
    string job_id = 1;
    int32 code = 2;              // Application error codes
    string message = 3;
}

message GetJobRequest { string job_id = 1; }
message GetJobResponse {
    enum Status { UNKNOWN = 0; QUEUED = 1; RUNNING = 2; SUCCEEDED = 3;
    FAILED = 4; }
    string job_id = 1;
    Status status = 2;
    string s3_url = 3;           // filled if SUCCEEDED
    string error = 4;            // filled if FAILED
    int32 progress = 5;         // 0..100
}

```

Rust (tonic) – service skeleton

```
use tonic::{Request, Response, Status};
use tonic::codegen::futures_core::Stream;
use tokio_stream::wrappers::ReceiverStream;

pub struct GeneratorSvc { /* state: model cache, redis, minio, etc. */ }

#[tonic::async_trait]
impl video::generator_server::Generator for GeneratorSvc {
    type GenerateStream = ReceiverStream<Result<video::GenerateEvent,
    Status>>;

    async fn Generate(
        &self,
        req: Request<video::GenerateRequest>,
    ) -> Result<Response<Self::GenerateStream>, Status> {
        // load or get cached model
        // spawn task producing Progress/Result/Error events
        // stream back via mpsc channel
        todo!()
    }

    async fn GetJob(
        &self,
        req: Request<video::GetJobRequest>,
    ) -> Result<Response<video::GetJobResponse>, Status> {
        // read from Redis/state store
        todo!()
    }
}
```

Python client (Celery task)

```
import grpc
from video_pb2 import GenerateRequest
from video_pb2_grpc import GeneratorStub

@celery.task(bind=True, acks_late=True)
def generate_task(self, job_id: str, prompt: str, options: dict):
    chan = grpc.insecure_channel("video-worker:50051") # or mTLS
    stub = GeneratorStub(chan)
    req = GenerateRequest(job_id=job_id, prompt=prompt, model="mochi-1-
    preview", scratch_dir="/scratch", options=options)

    for event in stub.Generate(req):
        if event.HasField("progress"):
            redis.hset(f"job:{job_id}", mapping={"status": "RUNNING",
            "progress": pct(event.progress)})
```

```

        elif event.HasField("result"):
            # upload may be done in Rust; optionally verify ETag
            redis.hset(f"job:{job_id}", mapping={"status": "SUCCEEDED",
"s3_url": event.result.s3_url})
        elif event.HasField("error"):
            redis.hset(f"job:{job_id}", mapping={"status": "FAILED",
"error": event.error.message})
            raise Exception(event.error.message)

```

API semantics:

- **Idempotency:** `job_id` is unique; repeated `Generate(job_id)` should resume or no-op.
- **Deadlines/timeouts:** client sets gRPC deadline; server respects and continues in background if desired.
- **Backpressure:** streaming progress avoids long-polling and supports cancellation.
- **Observability:** include step counts and GPU util in events for metrics.

Kubernetes

The examples assume namespace `video` and GPU scheduling via the NVIDIA Device Plugin/Operator.

Namespace & RBAC

```

apiVersion: v1
kind: Namespace
metadata:
  name: video
---
apiVersion: v1
kind: ServiceAccount
metadata:
  name: runtime
  namespace: video
---
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: runtime-role
  namespace: video
rules:
  - apiGroups: [""]
    resources: ["pods", "pods/log", "secrets", "configmaps"]
    verbs: ["get", "list", "watch"]
---
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: runtime-rb

```

```

namespace: video
subjects:
  - kind: ServiceAccount
    name: runtime
    namespace: video
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: Role
  name: runtime-role

```

Config & Secrets

```

apiVersion: v1
kind: Secret
metadata:
  name: s3-credentials
  namespace: video
stringData:
  AWS_ACCESS_KEY_ID: "minio"
  AWS_SECRET_ACCESS_KEY: "<redacted>"
  AWS_S3_ENDPOINT: "http://minio.video.svc.cluster.local:9000"
  AWS_S3_BUCKET: "videos"
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-config
  namespace: video
data:
  REDIS_URL: "redis://redis-master.video.svc.cluster.local:6379/0"
  SCRATCH_DIR: "/scratch"
  MODEL_NAME: "mochi-1-preview"

```

PVC (scratch)

```

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: video-scratch
  namespace: video
spec:
  accessModes: ["ReadWriteOnce"] # Use RWX if multiple pods read
  same path
  resources:
    requests:
      storage: 500Gi
  storageClassName: fast-nvme

```



```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: video-api
  namespace: video
spec:
  replicas: 2
  selector:
    matchLabels:
      app: video-api
  template:
    metadata:
      labels:
        app: video-api
      annotations:
        prometheus.io/scrape: "true"
        prometheus.io/port: "8000"
    spec:
      serviceAccountName: runtime
      containers:
        - name: api
          image: your-dockerhub/video-api:latest
          ports:
            - containerPort: 8000
          envFrom:
            - configMapRef: { name: app-config }
            - secretRef: { name: s3-credentials }
          readinessProbe:
            httpGet: { path: /healthz, port: 8000 }
            periodSeconds: 5
          livenessProbe:
            httpGet: { path: /livez, port: 8000 }
            periodSeconds: 10

```

```

apiVersion: v1
kind: Service
metadata:
  name: video-api
  namespace: video
spec:
  selector:
    app: video-api
  ports:
    - port: 80
      targetPort: 8000
      protocol: TCP

```

```

apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler

```

```

metadata:
  name: video-api
  namespace: video
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: video-api
  minReplicas: 2
  maxReplicas: 10
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 70

```

Rust Worker Deployment + Service + PDB

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: video-worker
  namespace: video
spec:
  replicas: 4
  selector:
    matchLabels:
      app: video-worker
  template:
    metadata:
      labels:
        app: video-worker
      annotations:
        prometheus.io/scrape: "true"
        prometheus.io/port: "9464" # if exposing metrics
    spec:
      serviceAccountName: runtime
      nodeSelector:
        nvidia.com/gpu.present: "true"
      tolerations:
        - key: nvidia.com/gpu
          operator: Exists
          effect: NoSchedule
      volumes:
        - name: scratch
          persistentVolumeClaim:
            claimName: video-scratch
      containers:

```

```

- name: worker
  image: your-dockerhub/video-worker:latest
  ports:
    - containerPort: 50051 # gRPC
  resources:
    limits:
      nvidia.com/gpu: 2
    requests:
      cpu: "2"
      memory: "8Gi"
  volumeMounts:
    - name: scratch
      mountPath: /scratch
  envFrom:
    - configMapRef: { name: app-config }
    - secretRef: { name: s3-credentials }
  readinessProbe:
    tcpSocket: { port: 50051 }
    periodSeconds: 5
  livenessProbe:
    tcpSocket: { port: 50051 }
    periodSeconds: 10
---
apiVersion: v1
kind: Service
metadata:
  name: video-worker
  namespace: video
spec:
  selector:
    app: video-worker
  ports:
    - port: 50051
      targetPort: 50051
      protocol: TCP
---
apiVersion: policy/v1
kind: PodDisruptionBudget
metadata:
  name: video-worker-pdb
  namespace: video
spec:
  minAvailable: 75%
  selector:
    matchLabels:
      app: video-worker

```

Network Policies (lock down traffic)

```

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: api-allow
  namespace: video
spec:
  podSelector:
    matchLabels: { app: video-api }
  ingress:
    - from:
        - podSelector: {} # adjust for ingress controller
      ports:
        - protocol: TCP
          port: 80
-----
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: worker-allow
  namespace: video
spec:
  podSelector:
    matchLabels: { app: video-worker }
  ingress:
    - from:
        - podSelector:
            matchLabels: { app: video-api }
        - podSelector:
            matchLabels: { app: celery } # if separate Celery
  deployment
    ports:
      - protocol: TCP
        port: 50051

```

Deployment Plan (cluster-level)

1. Verify cluster access

```
kubectl --kubeconfig kubeconfig.yml get nodes
```

2. Install NVIDIA GPU Operator + device plugin

```

helm repo add nvidia https://helm.ngc.nvidia.com/nvidia
helm repo update
helm install gpu-operator nvidia/gpu-operator -n gpu-operator --
create-namespace

```

3. Install Redis (Bitnami)

```
helm repo add bitnami https://charts.bitnami.com/bitnami
helm install redis bitnami/redis -n video --create-namespace \
  --set architecture=replication --set auth.enabled=false
```

4. Deploy API

```
kubectl apply -f k8s/api.yaml
```

5. Deploy Workers

```
kubectl apply -f k8s/worker.yaml
```

6. Create PVC

```
kubectl apply -f k8s/pvc.yaml
```

7. Deploy MinIO

```
helm install minio bitnami/minio -n video \
  --set resources.requests.memory=1Gi \
  --set mode=standalone \
  --set auth.rootUser=minio,auth.rootPassword=<redacted>
```

Scaling Guidelines

Scenario	API replicas	Worker pods	GPUs/pod	Storage Strategy
Many small jobs	4	4	2	PVC → MinIO
Heavy long jobs	2	4	2 (1 task/GPU)	PVC → MinIO
Mixed	HPA on RPS	Separate queues	1–2	PVC per worker
Very large outputs	2–3	4	2	Direct stream to MinIO

Notes:

- Consider **one task per GPU** to avoid context thrash.

- Use **separate Celery queues** (e.g., **short**, **long**) mapped to different workers.
- For multi-reader workflows, prefer **RWX** PVCs (e.g., NFS/CSI) or remove PVC and **stream directly** to MinIO.

API (FastAPI) – status & presigned URL

```
from fastapi import FastAPI, HTTPException
import redis, boto3, os

r = redis.Redis.from_url(os.environ["REDIS_URL"])
app = FastAPI()

@app.get("/jobs/{job_id}")
def get_job(job_id: str):
    data = r.hgetall(f"job:{job_id}") or {}
    if not data:
        raise HTTPException(404, "unknown job")
    return {k.decode(): v.decode() for k,v in data.items()}

@app.get("/download/{job_id}")
def presign(job_id: str):
    s3 = boto3.client("s3", endpoint_url=os.environ["AWS_S3_ENDPOINT"],
aws_access_key_id=os.environ["AWS_ACCESS_KEY_ID"],
aws_secret_access_key=os.environ["AWS_SECRET_ACCESS_KEY"])
    key = r.hget(f"job:{job_id}", "s3_url")
    if not key:
        raise HTTPException(404, "no artifact")
    bucket = os.environ["AWS_S3_BUCKET"]
    url = s3.generate_presigned_url("get_object", Params={"Bucket":
bucket, "Key": key.decode()}, ExpiresIn=3600)
    return {"url": url}
```

Security

- **mTLS** for gRPC between Celery and Rust workers (SPIRE/certs or cert-manager Issuer).
- **NetworkPolicy** to restrict gRPC to API/Celery only.
- **ServiceAccount** with minimal RBAC; avoid node-wide permissions.
- **Secrets**: use Kubernetes Secrets + optional KMS envelope encryption.
- **Pod Security**: run as non-root, drop **CAP_SYS_ADMIN**, read-only root FS where possible.

```
securityContext:
  runAsNonRoot: true
  runAsUser: 1000
  fsGroup: 1000
  allowPrivilegeEscalation: false
  readOnlyRootFilesystem: true
```

Observability

- **Metrics** (Prometheus): GPU utilization (DCGM exporter), queue depth, job throughput, success/failure rates, p95 latency, upload time.
- **Logs**: structured JSON logs from API/worker; forward via Fluent Bit or OpenTelemetry Collector.
- **Tracing**: OpenTelemetry SDK in API and Rust worker (OTLP → collector).
- **Alerts**: worker crashloop, PVC > 80% full, queue depth > threshold, high GPU idle.

Example Prometheus annotations are shown on Deployments above.

Development Workflow

1. Stand up local **FastAPI + Celery + Redis**.
 2. Integrate `**HuggingFace**`` (guard model weights with `.gitignore`; use `HF_TOKEN`).
 3. Test generation locally with GPU (Docker + `--gpus all`).
 4. Build Docker images for API & worker and push to registry.
 5. Apply manifests (`kubectl apply -f k8s/`), then `kubectl logs -n video -l app=video-worker`.
 6. Submit a job; verify progress and final artifact in MinIO.
-

Optional Frontend (MVP)

- Minimal SPA or HTML with:
 - Text prompt input
 - Submit button → POST `/jobs`
 - Status polling → GET `/jobs/{id}`
 - Download with presigned URL → GET `/download/{id}`
-

Hardening Checklist

-
-

Quick Commands

```
# Port-forward API
kubectl -n video port-forward svc/video-api 8080:80

# Tail worker logs
kubectl -n video logs -f deploy/video-worker -c worker

# Watch GPU pods
kubectl -n video get pods -l app=video-worker -w
```

